

Do LLM Rankings Change if You Ask Differently?

Ranking reversals under non-semantic prompt perturbations
and subset resampling

ST5230 Group 14

A Familiar Story, and the Question Behind It

What a leaderboard suggests

- Model A: 89.2%
- Model B: 87.3%
- Therefore, A is better than B

A Familiar Story, and the Question Behind It

What a leaderboard suggests

- Model A: 89.2%
- Model B: 87.3%
- Therefore, A is better than B

What we ask instead

- What if we ask the same question differently?
- What if we evaluate on a slightly different subset?
- Does the rank order remain the same?

A Familiar Story, and the Question Behind It

What a leaderboard suggests

- Model A: 89.2%
- Model B: 87.3%
- Therefore, A is better than B

What we ask instead

- What if we ask the same question differently?
- What if we evaluate on a slightly different subset?
- Does the rank order remain the same?

Our target is not only score accuracy, but the reliability of the ranking itself.

- Static leaderboards are attractive because they compress a benchmark into a single ordering.

Motivation

- Static leaderboards are attractive because they compress a benchmark into a single ordering.
- But that ordering implicitly assumes **evaluation stability**: small, reasonable protocol changes should not alter the winner.

Motivation

- Static leaderboards are attractive because they compress a benchmark into a single ordering.
- But that ordering implicitly assumes **evaluation stability**: small, reasonable protocol changes should not alter the winner.
- In practice, benchmark outcomes depend on at least two nontrivial factors:

Motivation

- Static leaderboards are attractive because they compress a benchmark into a single ordering.
- But that ordering implicitly assumes **evaluation stability**: small, reasonable protocol changes should not alter the winner.
- In practice, benchmark outcomes depend on at least two nontrivial factors:
 - prompt formulation

Motivation

- Static leaderboards are attractive because they compress a benchmark into a single ordering.
- But that ordering implicitly assumes **evaluation stability**: small, reasonable protocol changes should not alter the winner.
- In practice, benchmark outcomes depend on at least two nontrivial factors:
 - prompt formulation
 - evaluation subset composition

We study when a leaderboard is robust, and when it is only an artifact of one evaluation choice.

Motivation

- Static leaderboards are attractive because they compress a benchmark into a single ordering.
- But that ordering implicitly assumes **evaluation stability**: small, reasonable protocol changes should not alter the winner.
- In practice, benchmark outcomes depend on at least two nontrivial factors:
 - prompt formulation
 - evaluation subset composition
- If these sources of variation are large relative to model gaps, rank claims may be overconfident.

We study when a leaderboard is robust, and when it is only an artifact of one evaluation choice.

Benchmark papers

MMLU, ARC, and HellaSwag establish the task settings we evaluate.

Positioning Relative to Prior Work

Benchmark papers

MMLU, ARC, and HellaSwag establish the task settings we evaluate.

Prompt robustness

Prior work shows that task-preserving prompt changes can affect LLM behavior.

Positioning Relative to Prior Work

Benchmark papers

MMLU, ARC, and HellaSwag establish the task settings we evaluate.

Prompt robustness

Prior work shows that task-preserving prompt changes can affect LLM behavior.

Evaluation reliability

Recent discussions emphasize that protocol matters, not just model identity.

Positioning Relative to Prior Work

Benchmark papers

MMLU, ARC, and HellaSwag establish the task settings we evaluate.

Prompt robustness

Prior work shows that task-preserving prompt changes can affect LLM behavior.

Evaluation reliability

Recent discussions emphasize that protocol matters, not just model identity.

Our contribution

We focus on **ranking stability**: not just whether scores move, but whether the order between models reverses.

Research Questions

- 1 Do non-semantic prompt perturbations change model performance?

Research Questions

- ① Do non-semantic prompt perturbations change model performance?
- ② Does evaluation-subset variation change relative model rankings?

Research Questions

- ① Do non-semantic prompt perturbations change model performance?
- ② Does evaluation-subset variation change relative model rankings?
- ③ Are those changes large enough to create **ranking reversals**?

Research Questions

- ① Do non-semantic prompt perturbations change model performance?
- ② Does evaluation-subset variation change relative model rankings?
- ③ Are those changes large enough to create **ranking reversals**?

If the answer is no

A static leaderboard is a reasonable summary.

If the answer is yes

Rank claims should be reported with uncertainty and robustness diagnostics.

Method

Treat ranking as a random outcome induced by evaluation choices

Core Idea and Theoretical Framing

- Keep task semantics fixed.

Core Idea and Theoretical Framing

- Keep task semantics fixed.
- Perturb the prompt in a controlled, non-semantic way.

Core Idea and Theoretical Framing

- Keep task semantics fixed.
- Perturb the prompt in a controlled, non-semantic way.
- Resample the evaluation subset via bootstrap.

$r_{d,p,b}(m)$ = rank of model m

$$\text{RRR}_{d,p} = \Pr(r_{d,p,b} \neq r_{d,\text{minimal},b})$$

Core Idea and Theoretical Framing

- Keep task semantics fixed.
- Perturb the prompt in a controlled, non-semantic way.
- Resample the evaluation subset via bootstrap.
- Recompute the ranking under each induced evaluation configuration.

$r_{d,p,b}(m)$ = rank of model m

$RRR_{d,p} = \Pr(r_{d,p,b} \neq r_{d,\text{minimal},b})$

Interpretation

Rank is not treated as a fixed truth, but as a distribution over plausible evaluation choices.

Core Idea and Theoretical Framing

- Keep task semantics fixed.
- Perturb the prompt in a controlled, non-semantic way.
- Resample the evaluation subset via bootstrap.
- Recompute the ranking under each induced evaluation configuration.

$r_{d,p,b}(m)$ = rank of model m

$RRR_{d,p} = \Pr(r_{d,p,b} \neq r_{d,\text{minimal},b})$

Interpretation

Rank is not treated as a fixed truth, but as a distribution over plausible evaluation choices.

High RRR means a single leaderboard is a fragile snapshot rather than a stable conclusion.

Experimental Setup

Benchmarks

- MMLU: 300 questions
- ARC-Challenge: 300 questions
- HellaSwag: 300 questions

Prompt templates

- `minimal_instruction`
- `benchmark_style`
- `natural_style`
- `order_phrasing_variation`

Experimental Setup

Benchmarks

- MMLU: 300 questions
- ARC-Challenge: 300 questions
- HellaSwag: 300 questions

Prompt templates

- `minimal_instruction`
- `benchmark_style`
- `natural_style`
- `order_phrasing_variation`

Models

- GPT-4o-mini
- Gemini-2.0-flash
- Qwen-2.5-7B
- Llama-3.1-8B

Inference / analysis

- `temperature = 0`
- `max_tokens = 150`
- 2000 bootstrap resamples per dataset

Experimental Setup

Benchmarks

- MMLU: 300 questions
- ARC-Challenge: 300 questions
- HellaSwag: 300 questions

Prompt templates

- `minimal_instruction`
- `benchmark_style`
- `natural_style`
- `order_phrasing_variation`

Models

- GPT-4o-mini
- Gemini-2.0-flash
- Qwen-2.5-7B
- Llama-3.1-8B

Inference / analysis

- `temperature` = 0
- `max_tokens` = 150
- 2000 bootstrap resamples per dataset

900

questions

3600

prompt requests


14400

responses


3

parse failures

Pipeline



Pipeline




- This design isolates two realistic sources of instability: **prompt choice** and **subset choice**.
- It also keeps the task fixed, so reversals cannot be dismissed as semantic task drift.


Results

Three benchmarks, three very different stability stories

Main Result: Ranking Stability Is Benchmark-dependent




Main Result: Ranking Stability Is Benchmark-dependent



RRR relative to the minimal prompt

- ARC: 4.70%, 1.65%, 1.05%
- MMLU: 19.60%, 8.05%, 13.50%


Main Result: Ranking Stability Is Benchmark-dependent



RRR relative to the minimal prompt

- ARC: 4.70%, 1.65%, 1.05%
- MMLU: 19.60%, 8.05%, 13.50%
- HellaSwag: 62.55%, 41.35%, 38.50%


Main Result: Ranking Stability Is Benchmark-dependent



RRR relative to the minimal prompt

- ARC: 4.70%, 1.65%, 1.05%
- MMLU: 19.60%, 8.05%, 13.50%
- HellaSwag: 62.55%, 41.35%, 38.50%

Main Result: Ranking Stability Is Benchmark-dependent




RRR relative to the minimal prompt


- ARC: 4.70%, 1.65%, 1.05%
- MMLU: 19.60%, 8.05%, 13.50%
- HellaSwag: 62.55%, 41.35%, 38.50%

Same method, same models,
same prompt philosophy;
radically different stability.

Reading the Accuracy Landscape




Reading the Accuracy Landscape



- **ARC**: tight boxes, clear separation, wide performance margins.
- **HellaSwag**: larger prompt sensitivity, especially around the top two models.

Reading the Accuracy Landscape



- **ARC**: tight boxes, clear separation, wide performance margins.
- **HellaSwag**: larger prompt sensitivity, especially around the top two models.
- **MMLU**: clear winner, but less separation in the middle ranks.

Case Study 1: HellaSwag and the Fragile Winner

62.55%

highest RRR under benchmark-style
prompting

Case Study 1: HellaSwag and the Fragile Winner

62.55%

highest RRR under benchmark-style
prompting

1.92%

GPT lead over Gemini in
prompt-averaged accuracy

Case Study 1: HellaSwag and the Fragile Winner

62.55%

highest RRR under benchmark-style
prompting

1.92%

GPT lead over Gemini in
prompt-averaged accuracy

$p = 0.169$

the GPT–Gemini gap is not
significant

Case Study 1: HellaSwag and the Fragile Winner

62.55%

highest RRR under benchmark-style
prompting

1.92%

GPT lead over Gemini in
prompt-averaged accuracy

$p = 0.169$

the GPT–Gemini gap is not
significant

Interpretation

HellaSwag provides the clearest counterexample to the static leaderboard narrative: the apparent winner exists at the point-estimate level, but not as a robust rank claim.

Case Study 2: ARC-Challenge and a Robust Leaderboard

- Gemini is first under all prompt variants.

Case Study 2: ARC-Challenge and a Robust Leaderboard

- Gemini is first under all prompt variants.
- All pairwise comparisons are significant.

94.00%

Gemini prompt-averaged accuracy

Case Study 2: ARC-Challenge and a Robust Leaderboard

- Gemini is first under all prompt variants.
- All pairwise comparisons are significant.
- RRR remains close to zero across alternative prompts.

94.00%

Gemini prompt-averaged accuracy

99%+

Gemini top-1 probability across prompts

Case Study 2: ARC-Challenge and a Robust Leaderboard

- Gemini is first under all prompt variants.
- All pairwise comparisons are significant.
- RRR remains close to zero across alternative prompts.

ARC behaves like the idealized leaderboard story: prompt and subset variation are too weak to overturn the ordering.

94.00%

Gemini prompt-averaged accuracy

99%+

Gemini top-1 probability across prompts

1–5%

RRR range only

Case Study 3: MMLU and an Unstable Middle

- Gemini is clearly first and remains first under all prompt variants.

Case Study 3: MMLU and an Unstable Middle

- Gemini is clearly first and remains first under all prompt variants.
- But GPT vs Qwen is not statistically significant.

11.67%

Gemini lead over GPT

Case Study 3: MMLU and an Unstable Middle

- Gemini is clearly first and remains first under all prompt variants.
- But GPT vs Qwen is not statistically significant.
- So MMLU is not fully unstable; it is **hierarchically unstable**: the winner is settled, the middle ranks are not.

11.67%

Gemini lead over GPT

3.08%

GPT lead over Qwen

Case Study 3: MMLU and an Unstable Middle

- Gemini is clearly first and remains first under all prompt variants.
- But GPT vs Qwen is not statistically significant.
- So MMLU is not fully unstable; it is **hierarchically unstable**: the winner is settled, the middle ranks are not.

The right lesson is not “MMLU is noisy,” but “different parts of the ranking have different certainty levels.”

11.67%

Gemini lead over GPT


3.08%

GPT lead over Qwen

$p = 0.184$


GPT–Qwen gap is not significant

One More Layer: MMLU Is Also Subject-dependent




- Full MMLU preserves the same overall winner.

One More Layer: MMLU Is Also Subject-dependent




- Full MMLU preserves the same overall winner.
- But subject-level stability varies sharply.

One More Layer: MMLU Is Also Subject-dependent



- Full MMLU preserves the same overall winner.
- But subject-level stability varies sharply.
- Mean non-minimal RRR ranges from about 14.5% to 67.9%.

One More Layer: MMLU Is Also Subject-dependent



- Full MMLU preserves the same overall winner.
- But subject-level stability varies sharply.
- Mean non-minimal RRR ranges from about 14.5% to 67.9%.

Ranking robustness is benchmark-dependent **and** subdomain-dependent.

Takeaways

What this project says about how LLM evaluation should be reported

What We Learned

- ① Prompt perturbations and subset variation do not affect all benchmarks equally.

What We Learned

- ① Prompt perturbations and subset variation do not affect all benchmarks equally.
- ② Ranking reversals occur most often when performance gaps are narrow.

What We Learned

- ① Prompt perturbations and subset variation do not affect all benchmarks equally.
- ② Ranking reversals occur most often when performance gaps are narrow.
- ③ A single leaderboard may be robust, partially stable, or fragile depending on the task.

What We Learned

- ① Prompt perturbations and subset variation do not affect all benchmarks equally.
- ② Ranking reversals occur most often when performance gaps are narrow.
- ③ A single leaderboard may be robust, partially stable, or fragile depending on the task.
- ④ Average accuracy alone is not enough; rank uncertainty deserves explicit reporting.

What We Learned

- ① Prompt perturbations and subset variation do not affect all benchmarks equally.
- ② Ranking reversals occur most often when performance gaps are narrow.
- ③ A single leaderboard may be robust, partially stable, or fragile depending on the task.
- ④ Average accuracy alone is not enough; rank uncertainty deserves explicit reporting.

A better benchmark report should state not only who is best on average, but how robust that ordering is.

Implications and Limitations

Implications

- report confidence intervals and pairwise significance
- report ranking stability metrics such as RRR
- avoid overinterpreting narrow score gaps as stable superiority

Implications and Limitations

Implications

- report confidence intervals and pairwise significance
- report ranking stability metrics such as RRR
- avoid overinterpreting narrow score gaps as stable superiority

Limitations

- only 3 multiple-choice benchmarks
- only 4 models
- prompt perturbations are controlled and shallow, not exhaustive

Implications and Limitations

Implications

- report confidence intervals and pairwise significance
- report ranking stability metrics such as RRR
- avoid overinterpreting narrow score gaps as stable superiority

Limitations

- only 3 multiple-choice benchmarks
- only 4 models
- prompt perturbations are controlled and shallow, not exhaustive

This is not a claim that all leaderboards are unreliable. It is a claim that reliability must be demonstrated, not assumed.

ARC

robust leaderboard

Conclusion

ARC

robust leaderboard

MMLU

stable winner, unstable middle

Conclusion

ARC

robust leaderboard

MMLU

stable winner, unstable middle

HellaSwag

fragile top rank

ARC

robust leaderboard

MMLU

stable winner, unstable middle

HellaSwag


fragile top rank

The central question is not only “Who is best?” but “Does that ordering survive reasonable evaluation variation?”

Experiment Design and Result


Experimental workflow, metric definitions, and reversal cases we found

Pipeline




- Phase 1: samples **300 items** for each dataset and expands each item into **4 prompt templates**. For MMLU, we selected **6 distinct subjects** and sampled **50 items** from each subject.

Pipeline



- Phase 1: samples **300 items** for each dataset and expands each item into **4 prompt templates**. For MMLU, we selected **6 distinct subjects** and sampled **50 items** from each subject.
- Phase 2: queries 4 models on all prompt requests, parses each of the 14,400 responses into A/B/C/D, and turn it into a binary score.

Pipeline



- Phase 1: samples **300 items** for each dataset and expands each item into **4 prompt templates**. For MMLU, we selected **6 distinct subjects** and sampled **50 items** from each subject.
- Phase 2: queries 4 models on all prompt requests, parses each of the 14,400 responses into A/B/C/D, and turn it into a binary score.
- Phase 3: By bootstrapping, compute model accuracy and ranking reversal rates (see next slide).

Definition of Core Statistics: Accuracy and RRR's

- **Accuracy:** for any item set A and template T , a model's accuracy on A under T is the mean of its fixed binary scores when asked problems from A with template T .

Definition of Core Statistics: Accuracy and RRR's

- **Accuracy:** for any item set A and template T , a model's accuracy on A under T is the mean of its fixed binary scores when asked problems from A with template T .
- **Ranking:** for any item set A and template T , the 4 models can be sorted by their accuracy on A under T , inducing a ranking, denoted $R_{A,T}$.

Definition of Core Statistics: Accuracy and RRR's

- **Accuracy**: for any item set A and template T , a model's accuracy on A under T is the mean of its fixed binary scores when asked problems from A with template T .
- **Ranking**: for any item set A and template T , the 4 models can be sorted by their accuracy on A under T , inducing a ranking, denoted $R_{A,T}$.
- **Prompt-wise Ranking Reversal Rate (PRRR)**: For an item set P , two templates T_1 and T_2 , draw N subsets $P_i \subset P$; the PRRR on P between T_1 and T_2 is the frequency of different rankings occurring under two templates among the N subsets, denote $\text{PRRR}(P, T_1, T_2)$. i.e.,

$$\text{PRRR}(P, T_1, T_2) := \#\{i \in [1 : N] : R_{P_i, T_1} \neq R_{P_i, T_2}\} / N.$$

Definition of Core Statistics: Accuracy and RRR's

- **Accuracy:** for any item set A and template T , a model's accuracy on A under T is the mean of its fixed binary scores when asked problems from A with template T .
- **Ranking:** for any item set A and template T , the 4 models can be sorted by their accuracy on A under T , inducing a ranking, denoted $R_{A,T}$.
- **Prompt-wise Ranking Reversal Rate (PRRR):** For an item set P , two templates T_1 and T_2 , draw N subsets $P_i \subset P$; the PRRR on P between T_1 and T_2 is the frequency of different rankings occurring under two templates among the N subsets, denote $\text{PRRR}(P, T_1, T_2)$. i.e.,

$$\text{PRRR}(P, T_1, T_2) := \#\{i \in [1 : N] : R_{P_i, T_1} \neq R_{P_i, T_2}\} / N.$$

- **Subset Ranking Reversal Rate (SRRR):** For an item set P and template T , draw N subsets $P_i \subset P$; the SRRR on P under T is the frequency of different rankings occurring under T when asked subsets and the full set, denote $\text{SRRR}(P, T)$. i.e.,

$$\text{SRRR}(P, T) := \#\{i \in [1 : N] : R_{P, T} \neq R_{P_i, T}\} / N.$$

Definition of Core Statistics: Accuracy and RRR's

**We isolate two sources of instability: prompt choice and incomplete test set.
PRRR measures prompt-to-prompt rank disagreement on the same task set;
SRRR measures full-set vs subset rank instability under the same prompt.**


Definition of Core Statistics: Accuracy and RRR's

We isolate two sources of instability: prompt choice and incomplete test set. PRRR measures prompt-to-prompt rank disagreement on the same task set; SRRR measures full-set vs subset rank instability under the same prompt.

Experiment settings:

- We applied bootstrap number $N = 1000$.
- For 3 datasets (MMLU, ARC-Challenge, and HellaSwag), and 6 subjects of question sets from MMLU, we computed:
 - PRRR between minimal instruction and the other 3 templates;
 - SRRR under all 4 templates.


Experiment Results: Dataset-Level



PRRR Results:

- ARC-Challenge:
max PRRR = 10.6%;
- HellaSwag:
max PRRR = 56.0%;
- MMLU:
max PRRR = 15.6%.

Experiment Results: Dataset-Level



SRRR Results:

- ARC-Challenge:
max SRRR = 11.2%
- HellaSwag:
max SRRR = 42.8%
- MMLU:
max SRRR = 17.5%

Experiment Results: Dataset-Level


- **HellaSwag**: strongest dataset-level reversal case, especially under T2 Benchmark.
- **MMLU**: overall winner is stable, but middle ranks are more sensitive to subset changes.
- **ARC**: rank order survives both prompt and subset variation in most resamples.

Experiment Results: Dataset-Level

- **HellaSwag**: strongest dataset-level reversal case, especially under T2 Benchmark.
- **MMLU**: overall winner is stable, but middle ranks are more sensitive to subset changes.
- **ARC**: rank order survives both prompt and subset variation in most resamples.


HellaSwag is the clearest fragile-leaderboard case. ARC-Challenge remains the most stable benchmark.

Appendix V: MMLU Subject-Level Cases




- Highest subject PRRR: philosophy at 79.50%

Appendix V: MMLU Subject-Level Cases




- Highest subject PRRR: philosophy at 79.50%
- Highest subject SRRR: international_law at 71.30%

Appendix V: MMLU Subject-Level Cases



- Highest subject PRRR: philosophy at 79.50%
- Highest subject SRRR: international_law at 71.30%
- high_school_macro_economics is the most stable subject slice

Appendix V: MMLU Subject-Level Cases



- Highest subject PRRR: philosophy at 79.50%
- Highest subject SRRR: international_law at 71.30%
- high_school_macro_economics is the most stable subject slice

MMLU's aggregate stability hides strong subject-level heterogeneity.